

```
import torch.nn as nn
import torch.utils.model_zoo as model_zoo
import torch
import torchvision.datasets as datasets
import torchvision.transforms as transforms
import torch.optim as optim

def conv3x3(in_planes, out_planes, stride=1):
    """3x3 convolution with padding"""
    return nn.Conv2d(in_planes, out_planes, kernel_size=3, stride=stride, padding=1, bias=False)

def conv1x1(in_planes, out_planes, stride=1):
    """1x1 convolution"""
    return nn.Conv2d(in_planes, out_planes, kernel_size=1, stride=stride, bias=False)

class BasicBlock(nn.Module):
    expansion = 1

    def __init__(self, inplanes, planes, stride=1, downsample=None):
        super(BasicBlock, self).__init__()
        self.conv1 = conv3x3(inplanes, planes, stride)
        self.bn1 = nn.BatchNorm2d(planes)
        self.relu = nn.ReLU(inplace=True)
        self.conv2 = conv3x3(planes, planes)
        self.bn2 = nn.BatchNorm2d(planes)
        self.downsample = downsample
        self.stride = stride

    def forward(self, x):
        identity = x

        out = self.conv1(x)
        out = self.bn1(out)
        out = self.relu(out)
```

```

out = self.conv2(out)
out = self.bn2(out)

if self.downsample is not None:
    identity = self.downsample(x)

out += identity
out = self.relu(out)

return out

```

```
class CifarResNet(nn.Module):
```

```

def __init__(self, block, layers, num_classes=100):
    super(CifarResNet, self).__init__()
    self.inplanes = 16
    self.conv1 = conv3x3(3, 16)
    self.bn1 = nn.BatchNorm2d(16)
    self.relu = nn.ReLU(inplace=True)

    self.layer1 = self._make_layer(block, 16, layers[0])
    self.layer2 = self._make_layer(block, 32, layers[1], stride=2)
    self.layer3 = self._make_layer(block, 64, layers[2], stride=2)

    self.avgpool = nn.AdaptiveAvgPool2d((1, 1))
    self.fc = nn.Linear(64 * block.expansion, num_classes)

    for m in self.modules():
        if isinstance(m, nn.Conv2d):
            nn.init.kaiming_normal_(m.weight, mode='fan_out', nonlinearity='relu')
        elif isinstance(m, nn.BatchNorm2d):
            nn.init.constant_(m.weight, 1)
            nn.init.constant_(m.bias, 0)

    def _make_layer(self, block, planes, blocks, stride=1):
        downsample = None
        if stride != 1 or self.inplanes != planes * block.expansion:
            downsample = nn.Sequential(

```

```

        conv1x1(self.inplanes, planes * block.expansion, stride),
        nn.BatchNorm2d(planes * block.expansion),
    )

    layers = []
    layers.append(block(self.inplanes, planes, stride, downsample))
    self.inplanes = planes * block.expansion
    for _ in range(1, blocks):
        layers.append(block(self.inplanes, planes))

    return nn.Sequential(*layers)

def forward(self, x):
    x = self.conv1(x)
    x = self.bn1(x)
    x = self.relu(x)

    x = self.layer1(x)
    x = self.layer2(x)
    x = self.layer3(x)

    x = self.avgpool(x)
    x = x.view(x.size(0), -1)
    x = self.fc(x)

    return x

# Code added : To Tune hyper-parameters
_num_epoch = 42
_lr = 0.001
_momentum = 0.9
_weight_decay = 1e-4

# Code added to update learning rate to decay lr after each epoch
def update_lr(optimizer, lr):
    for param_group in optimizer.param_groups:
        param_group['lr'] = lr

```

```

class cifar_resnet20(nn.Module):
    def __init__(self):
        super(cifar_resnet20, self).__init__()
        ResNet20 = CifarResNet(BasicBlock, [3, 3, 3])
        url = 'https://github.com/chenyaofo/pytorch-cifar-models/releases/download/resnet/cifar100_resnet20-23dac2f

        ResNet20.load_state_dict(model_zoo.load_url(url))
        modules = list(ResNet20.children())[:-1]
        backbone = nn.Sequential(*modules)
        self.backbone = nn.Sequential(*modules)
        self.fc = nn.Linear(64, 10)

    def forward(self, x):
        out = self.backbone(x)
        out = out.view(out.shape[0], -1)
        out = self.fc(out)
        return out

if __name__ == '__main__':
    model = cifar_resnet20().cuda()

    transform = transforms.Compose([transforms.ToTensor(), transforms.Normalize(mean=(0.499, 0.499, 0.499), std=(0.1
    trainset = datasets.CIFAR10('./data', download=True, transform=transform)
    trainloader = torch.utils.data.DataLoader(trainset, batch_size=32, shuffle=True, num_workers=2)

    # Code added : To create testset and testloader
    testset = datasets.CIFAR10(root='./data', train=False, download=True, transform=transform)
    testloader = torch.utils.data.DataLoader(testset, batch_size=4, shuffle=False, num_workers=2)

    criterion = nn.CrossEntropyLoss()

    # Code added : To add L2 regularizer by setting weight_decay
    optimizer = optim.SGD(list(model.fc.parameters()), lr = _lr, momentum = _momentum, weight_decay = _weight_decay

    ## Do the training
    for epoch in range(_num_epoch): # loop over the dataset multiple times
        running_loss = 0.0

```

```

for i, data in enumerate(trainloader, 0):
    # get the inputs
    inputs, labels = data
    # zero the parameter gradients
    optimizer.zero_grad()

    # forward + backward + optimize
    outputs = model(inputs.cuda())
    loss = criterion(outputs, labels.cuda())
    loss.backward()
    optimizer.step()
    running_loss += loss.item()
    if i % 1000 == 999:    # print every 1000 mini-batches
        print('[%d, %5d] loss: %.3f' %
              (epoch + 1, i + 1, running_loss / 1000))
        running_loss = 0.0

# Decay learning rate
if (epoch+1) % 20 == 0:
    _lr /= 3
    update_lr(optimizer, _lr)

print('==== Finished Training =====')

# Code added : To calculate the accuracy rate on test error and to save best model
print('Testing 10000 images for : ' + str(epoch + 1) + " epoch(s)")
best_acc = 0
correct = 0
total = 0
with torch.no_grad():
    for i, data in enumerate(testloader, 0):
        images, labels = data
        outputs = model(images.cuda())
        _, predicted = torch.max(outputs, 1)
        total += labels.size(0)
        correct += (predicted.cuda() == labels.cuda()).sum().item()
print('Accuracy : ' + str((100 * correct / total)))

```

```
if(best_acc < (100 * correct / total)):
    best_acc = (100 * correct / total)
    torch.save(model, 'model_best.pth')
    print("New Best Model Saved")
```

Downloading: "https://github.com/chenyafo/pytorch-cifar-models/releases/download/resnet/cifar100_resnet20-23
100% 1.11M/1.11M [00:00<00:00, 3.51MB/s]

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to ./data/cifar-10-python.tar.gz
170499072/? [00:04<00:00, 57689996.77it/s]

Extracting ./data/cifar-10-python.tar.gz to ./data
Files already downloaded and verified

```
[1, 1000] loss: 1.313
[2, 1000] loss: 1.013
[3, 1000] loss: 0.976
[4, 1000] loss: 0.974
[5, 1000] loss: 0.967
[6, 1000] loss: 0.966
[7, 1000] loss: 0.956
[8, 1000] loss: 0.955
[9, 1000] loss: 0.961
[10, 1000] loss: 0.955
[11, 1000] loss: 0.955
[12, 1000] loss: 0.959
[13, 1000] loss: 0.952
[14, 1000] loss: 0.960
[15, 1000] loss: 0.950
[16, 1000] loss: 0.947
[17, 1000] loss: 0.947
[18, 1000] loss: 0.947
[19, 1000] loss: 0.958
[20, 1000] loss: 0.955
[21, 1000] loss: 0.944
[22, 1000] loss: 0.941
[23, 1000] loss: 0.942
[24, 1000] loss: 0.942
[25, 1000] loss: 0.952
[26, 1000] loss: 0.943
[27, 1000] loss: 0.945
[28, 1000] loss: 0.943
[29, 1000] loss: 0.944
[30, 1000] loss: 0.942
[31, 1000] loss: 0.948
[32, 1000] loss: 0.946
[33, 1000] loss: 0.945
[34, 1000] loss: 0.950
[35, 1000] loss: 0.948
```

✓ 42m 13s completed at 22:43

